

KANTIPUDI KESAVA SAI VEERENDRA

RAJAHMUNDRY | kantipudikesavasaiveerendra@gmail.com | 8790418895 | kesavakantipudi.vercel.app |
linkedin.com/in/kesavakantipudi | github.com/kesavakantipudi

SUMMARY

AI/ML Engineering student with hands-on experience building LLM applications, RAG systems, scalable AI APIs, and cloud-native solutions. Strong foundation in data structures, algorithms, machine learning, and distributed systems, with 1000+ coding problems solved across competitive programming platforms. Passionate about designing production-ready AI systems that scale to real-world workloads.

Technical Skills

Programming: Python, SQL, Java
Computer Science: Data Structures & Algorithms, OOP, DBMS, Operating Systems
Machine Learning: Classification, Regression, Clustering, Feature Engineering, Model Evaluation
Deep Learning & LLM: PyTorch, CNN, Transformers, LoRA Fine-Tuning
GenAI Systems: RAG Pipelines, Vector Databases, Embeddings, AI Safety, Guardrails, Prompt & Context Engineering
Backend & Deployment: FastAPI, REST APIs, Asynchronous Processing, Docker
Data & Tools: NumPy, Pandas, Scikit-learn, HuggingFace, Redis, PostgreSQL
Cloud: AWS, Google Cloud Platform, Git, Linux

Education

B.Tech - AI & ML, Aditya College of Engineering & Technology 2023 - 2027 | GPA: 8.5/10
Intermediate (MPC), Aditya Junior College 2021 - 2023 | 89.1%

Experience

Data Specialist Intern - Technical Hub May 2025 - June 2025

- o Built data processing and analytics workflows for educational datasets.
- o Automated reporting pipelines, reducing manual effort and improving reporting efficiency.
- o Analyzed large datasets to generate actionable insights supporting operational decisions.

Projects

Event-Driven RAG Document Ingestion Pipeline | Python, RabbitMQ, ChromaDB, Docker | [GITHUB](#)

- o Built an event-driven document ingestion pipeline using RabbitMQ and vector databases for scalable knowledge processing.
- o Implemented asynchronous document processing, chunking, embedding generation, and indexing workflows.
- o Designed a modular microservice architecture enabling independent scaling of ingestion and retrieval components.
- o Containerized services using Docker to support reproducible deployments and distributed execution.

Scalable AI Inference Service | Python, FastAPI, Redis, Docker | [GITHUB](#)

- o Developed a production-style AI inference API supporting asynchronous request processing and concurrent workloads.
- o Implemented Redis-based caching and rate limiting to optimize resource utilization and reduce inference latency.
- o Designed scalable request handling architecture separating API serving and model execution layers.
- o Built REST endpoints for image analysis and prediction workflows following production engineering practices.

Collaborative AI Agent Orchestration Framework with LangGraph | Python, LangGraph, LLMs | [GITHUB](#)

- o Developed a multi-agent AI framework capable of coordinating specialized agents to solve complex workflows.
- o Designed agent communication and orchestration pipelines for task planning, execution, and result aggregation.
- o Implemented modular agent architecture supporting extensibility and reusable workflow components.
- o Explored agentic AI patterns including routing, memory management, and collaborative decision making.

Certifications

AWS Certified AI Practitioner | [\(Certificate\)](#)
SnowProAssociate: Platform | [\(Certificate\)](#)
Oracle: Java Certified Foundations Associate | [\(Certificate\)](#)
Oracle: Certified Foundations Associate - Database | [\(Certificate\)](#)
Microsoft Certified: Power Platform Developer Associate | [\(Certificate\)](#)
Microsoft Certified: Power Platform Solution Architect Expert | [\(Certificate\)](#)
Certiport: IT Specialist - HTML and CSS | [\(Certificate\)](#)

Coding

- o Solved 1000+ coding problems across **LeetCode, CodeChef, GeeksforGeeks, and HackerRank**
- o Strong foundation in Data Structures & Algorithms (DSA), Competitive Programming, and Problem Solving.